

Sampling Based Motion Planning Algorithms

299r Fall 2017 by Nao Ouyang

December 14, 2017

Abstract

This paper covers the work I did for my research rotation with Professor Lucas Janson in Fall of 2017. The rotation focused around sampling-based motion planning algorithms which are frequently used in robotics (my area of interest). The rotation can be roughly divided into three main components: 1. Understanding the algorithms 2. Understanding the proofs of the algorithmic properties 3. Computer simulations of the algorithms. The specific algorithms covered include: 1. RRT and it's optimal variant, RRT* 2. PRM and it's simplified variant, sPRM 3. FMT* Finally, I'll write a section reflecting on my experiences this semester.

1 Introduction

The goal of sampling-based algorithms is to reduce high-dimensionality problems, which in practice cannot be solved exactly, to something manageable by sampling randomly within the space to find paths.

The two foundational algorithms are Rapidly Exploring Random Trees, aka RRT; and Probabilistic Roadmaps, aka PRM. The first is used for on-line path planning, as it rapidly converges to a feasible path with increasing number of samples. On the other hand, probabilistic roadmaps make a graph covering the space. This graph can be reused for multiple start and goal locations, whereas the RRT algorithm generates a path for a specific pair of start and end points.

Generally proofs cover the feasibility, optimality, and computational complexity of the different algorithms. Researchers seek to characterize properties that help determine how many samples might be required to cover a space well enough to have a high probability of finding a feasible path (if it exists), including in terms of how fast this probability changes as (for instance) a function of the number of samples.

Although here we will present 2D (xy) examples of the algorithms where the robot is a point, these algorithms can be applied to any dimensions. For instance, we can imagine the dimensions for a multirotor might include x,y,z and also roll, pitch, and yaw of the copter. However, as the number of dimensions we must sample in increases, we encounter what is called the "curse of dimensionality". In many of the proofs, one will encounter a $1/dim$ exponent where, for instance, the convergence of a solution will slow down drastically the number of dimensions increases.

1.1 How people represent these problems

In all of these algorithms, we have a perfect map of the environment and can check for node and edge collisions. Additionally, we have a goal region rather than a point goal, as the probability of sampling exactly the goal node is essentially zero.

!! TODO: go over standard problem formulation, e.g. X_{free} , $B(\mu, r)$, etc.

2 Algorithms

Both PRM and RRT algorithms are straightforward, requiring less than 10 lines of pseudocode each ¹

The pseudocode for the algorithms may be found in [2]. Here I will attempt a more casual description.

!! TODO: go over all the symbols used e.g. set notation

2.1 PRM

We have two variables we control: the number of samples we use, which is pre-determined, and also the connection radius. In later cases we will write this connection radius as a function of the number of samples used in a given run.

The general idea of PRM is to sample uniformly over the free space (generally, we actually *close* to uniformly by sampling over the entire space, collision checking, and throwing out colliding nodes). We then have a "local planner" which connects each point to its neighbors. Two commonly used methods are "k-nearest", in which points are connected to its k nearest neighbors; and otherwise by a distance threshold. (for the purposes of this report, we will work in metric space, where the triangle inequality holds, and use Euclidean distance as the distance metric). If we sample enough points, according to percolation theory, we will end up with a well-connected graph (as opposed to isolated clusters).

Finally, when given a specific set of start and end points, we can use any standard graph search algorithm, such as A*, to find a path in the graph between the start and end points.

Simplified PRM, or sPRM, differs only in that connections between vertices in already connected components are allowed. This makes the implementation easier. In the original PRM, these connections were avoided as they do not contribute as much to the connectivity of the graph.

2.2 RRT

In RRT, we build a tree instead of a graph. This tree grows from the start node and terminates when it connects to the goal node. Whereas in PRM where we sample all the points at once and then run a local planner and finally a graph search, in RRT we iteratively randomly sample a point in space. In this way RRT is more suitable for online applications, where we care more about finding a solution fast than have a good coverage of the entire space.

We first include the starting point in the tree. Then, we randomly sample a point in free space, and then connect it to the nearest node in the tree. If the neighbor is outside the connection radius, then we "steer" (or take a step of distance r from the neighbor to the sample), and add that closer node to the tree (instead of the original sampled point). Of course, the edge we add must be collision free.

Of note, randomly sampling in the entire space is preferable to, for instance, sampling random control inputs from the existing tree. In the latter case, new points will tend to cluster around existing points. With sampling across all of free space, points will tend to land in unexplored space (since the unexplored space will be of greater area than the explored space, thus the probability of sampling in unexplored space is bigger). Voronoi diagrams can help visualize this effect of random sampling.

!! Todo: add picture of voronoi diagrams

¹Pseudocode as written in that "algorithm" format all CS papers seem to use, with the \leftarrow signs.

2.3 PRM*

In order to guarantee that PRM will be asymptotically optimal, we simply define a lower limit on the connection radius as a function of the number of samples.

2.4 RRT*

An unfortunate downside of RRT is that, because it stops growing the tree as soon as it reaches the goal node, it is likely to return a suboptimal route. In fact, it can be proven that as $\lim_{n \rightarrow \infty}$, RRT is **guaranteed** to return a suboptimal path.

In order to guarantee that RRT will return an optimal path as $\lim_{n \rightarrow \infty}$, we must introduce a "rewire" function. For each point we put down, we will rewire it and its neighbors to make sure that the each point in the tree is reached with least cost (all paths are "shortest paths"). Rewiring involves removing an edge and replacing it with another edge. Additionally, as the number of nodes we have sampled increases, we decrease the connection radius (as a function of n).

This rewiring step is done for two conditions.

Specifically, when we add a node, we check that the edge to its nearest neighbor in the tree is collision-free (that is, we check that the node can be added to the tree at all). After that, we consider the edges to all points within a radius of the new point. If one of edges results in a lower path cost (where path cost is defined starting from the origin) than the edge we just created, we erase the first edge and add in the latter.

After that, we try all the edges to neighbors in distance r from the new point, and if that new path to the neighboring edge results in a smaller cost than the current path to the neighboring edge (from the origin), then we deleted the parent and add in the edge connecting x_{new} to the neighbor.

!! Todo: image of example of rewiring step's two cases

2.5 FMT* (Fast Marching Tree)

2.5.1 Overview

This algorithm works some black magic which I assume is inspired by or building on some other method, because otherwise it's a pretty weird combo of things to come up with and then have great faith to put effort into proving its optimality. ²

The algorithm proceeds roughly (ignoring set considerations) as follows. Similar to PRM, we sample all n points at once. Then, we begin the tree at the start point. We look from a given existing node in that tree From there to its neighbors, as defined by a limit on the euclidean distance (cost ball). Within all these $m = \text{neighbors}(x, r)$ we find the lowest cost node, call it z . Then find z 's neighbors, which we might denote $y = \text{neighbors}(z, r)$. We look from to see if any of these y nodes connect with x in a collision free manner.

From the graph perspective, we are locating the proper node via a sort of heapsort technique ³ and then keeping track of, and growing, the set of vertices in a way that we do not re-calculate any nodes we have dealt with already. ^{4 5}

²Actually, how does that work? That sounds incredibly risky (sinking time into a proof of a property that may or may not exist) and/or a very in-the-dark way of finding new algorithms (come up with some idea and run simulations to see if its likely to have asymptotic optimality?

³by sort of checking the intersection of balls between neighbors of sampled points and neighbors of a node in the tree?

⁴I'm 100% guessing as to if this an hand-wavy intuitive way of understanding this algorithm and its advantages

⁵More formally, according to the paper [3] we are using ideas from the Fast Marching Method, where we "use a heapsort technique to located the proper sample point to update, and incrementally build the solution in an 'outward' direction, so that the algorithm does not need to backtrack over previously evaluation sample points." This 'one-pass property is what makes both the FMM and FMT* (in addition to its lazy strategy) particularly efficient.

2.5.2 In more detail

We maintain three sets of vertices: V_{open} , $V_{unvisited}$, and V_{closed} .

Note: I don't really have a better way of explaining it than what the paper says, so skipping this section for now. Refer to [3].

3 Theoretical Studies

In this section I've give a brief high-level idea of how the proofs of PRM, RRT, PRM*, and FMT* go about being, well, proofs.

In particular, I worked through proofs on

- probabilistic completeness, the likelihood as n goes to infinity that a feasible path will be found if it exists. since we are sampling randomly, we define probabilistic completeness instead of completeness
- asymptotics optimality, whether as n goes to infinity the pathcost asymptotically approaches the optimal pathcost

All the proofs use the idea of "distance balls" or "cost balls".

As show in Table 1 of [2]:

Algorithm	Probabilistic Completeness	Asymptotic Optimality
PRM	Yes	No
sPRM	Yes	Yes
RRT	Yes	No
PRM*	Yes	Yes
RRT*	Yes	Yes
FMT*	Yes	Yes

One of the important things is rate of convergence vs. number of samples or even more useful, change in path cost vs. time.

Of note, the proofs of sPRM and RRT show how they converge to optimality with at something like n^2

3.1 Probabilistic Completeness

3.1.1 PRM

The formal version of this proof is covered in [1]. ⁶

This paper gave a theoretical analysis of PRM. Specifically, it studied the 'dependence of failure probability' of finding a feasible path (if it exists) as a function of

- Number of nodes N in the roadmap.
- Distance from obstacles (nearest or on average)
- The length of the path

A use case is providing an estimate of the number N in order to achieve some desired bound on the failure rate.

The analysis is done on sPRM, the simplified version of PRM.

A casual explanation of the proof of the first theorem is as follows:

⁶This is the only proof I feel like I have a firm grasp on.

$$P(Failure) \leq \frac{2L}{R} \left(1 - \frac{\pi R^2}{4|F|}\right)^N \quad (1)$$

Given a path γ of length L , as well as the shortest distance between any point on the path and any obstacle, R : We place theoretical points equally spaced at $\frac{R}{2}$ apart (arclength, that is distance along the line and not euclidean distance, because it makes more sense to do so) on the path (if n is the number of of such points, $n = \frac{2L}{R}$). Then draw balls of radius R around each point. The area covered by these balls is guaranteed to be collision free, by definition of R . We then want to write the probability of failure as a function of the likelihood of at least one point being sampled within each of these balls, so that we may connect them to form a feasible path (does not lie within obstacles). However, we have to consider that a line connecting two such points in these space balls might cross the "dip" between any two balls. Thus, in order to make this proof work, we define further a set of balls, centered at the same points as before, but now with half the radius $- R/2$. The diameter of these smaller balls (B_{small}) is thus R . We can guarantee that so long as any two points fall within the area covered by the B_{small} s, we will be able to connect them in a straight line without leaving the collision free area as defined by the \cup total area covered by the B_{big} balls.

The maximum distance two points can lie apart from each other is thus $\frac{3R}{2}$.

We can now write an upper bound on the probability of failure. The probability of failure will always (strictly) be less than the likelihood of not picking a sample in a ball, times n times.

The above statement uses Boole's theorem, which in essence states that the probability of some events happening within the union of some set of independent events, will always be \leq the sum of the probabilities of each individual event happening.

$$P\left(\bigcup_{i=1}^n A_i\right) \leq \sum_{i=1}^n P(A_i) \quad (2)$$

Theorem 1 is given as:

$$\begin{aligned} P(failure) &\leq P(\text{some ball is empty}) \\ &\leq \sum_{j=1}^{n-1} P(\text{The } j_{th} \text{ ball is empty}) \\ &= \left(\lceil \frac{2L}{R} \rceil - 1\right) \left(1 - \frac{B_{R/2}}{|F|}\right)^N \end{aligned} \quad (3)$$

Since we have uniform sampling, and the each sample is taken independently, in 2D we can relate sampling probabilities as areas. For any given ball, the probability of not sampling within a given ball equals the the probability of picking a sample in free space but not within the ball. Since in 2D the area of the small balls is $\pi R^2/4$, the sum of this happening z times is

$$z \left(1 - \frac{\pi R^2}{4|F|}\right)^z \quad (4)$$

We calculate the likelihood of this happening for every single one of the n balls, minus the area of the half small ball covering a path behind the start point and the half small ball covering the area ahead of the final point. As these are independent events, we get

$$P(failure) \leq \left(\frac{2L}{R}\right) \left(1 - \frac{\pi R^2}{4|F|}\right)^N \quad (5)$$




Figure 1: FMT ball diagram

3.1.2 BREAK: REST IS TODO

As usual... I finish the PRM algorithm out of 5 or 6 and run out of time for the rest... Will return to this after 5pm 14 Dec 2017. I feel like I only really learning about PRM and reasons grumble about Julia

3.1.3 todo: the rest: rrt fmt* rrt* and all of AO

3.1.4 RRT

RRT proof [2] introduces notion of "clearance", strong and weak clearance. Basically finding an exact path is very hard to work with, but if we have an area, we can start talking about the likelihood of finding a path in that area.

Briefly, the idea is that, as we sample uniformly outside of the space, we will uniformly expand the boundary of space that the tree can reach into in one step (bound by the connection radius). Eventually this frontier will grow to encompass the whole space.

3.1.5 FMT*

we can think about covering the path with balls as usual.

Tight approximation to most of the path vs. loose approximation: The idea is to shrink the size of the balls so that the cost approaches optimal... or something... I can't interpret lemma 4.2 at the moment

I will probably write this up next, there is a lot of interesting translations into english that I learned.

3.2 AO

TODO <> I am not sure I really got around to **any** of the AO proofs.

3.3 computational complexity

3.3.1 computational complexity

Generally speaking, one of the biggest bottlenecks is collision checking.

Whatever we can do to minimize the amount of collision checking or speed it up will help the most.

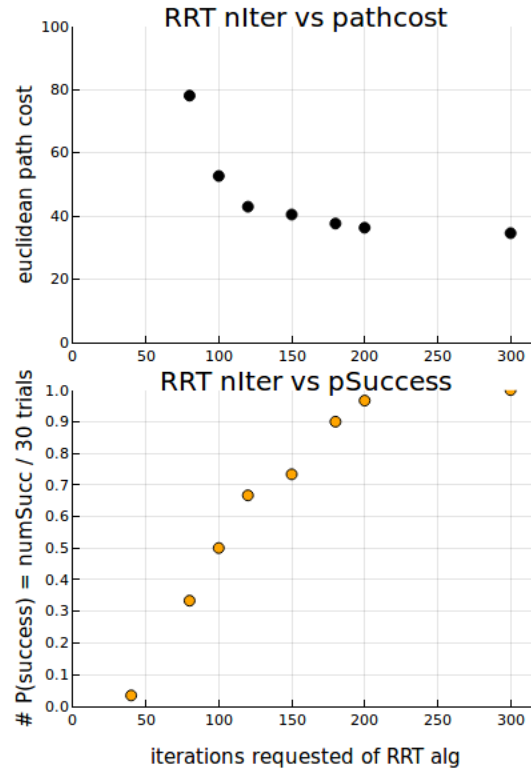


Figure 2: RRT

Additionally, caching graphs in for instance PRM, as well as intelligently determining places to sample, if a path is not found, could be important.

I'm not sure what to say other than that... collision checking seems obvious.

Of course, in my code the nearest neighbor lookups also took a lot of time, but that is because I implemented it in the simplest way possible.

Perhaps if at some point with large enough dimension, the edge and vertice graph may have to be stored in a database. In this case read and write speeds may become important (minimizing calls), or perhaps the wireless connection bandwidth. But I don't think this would ever be the bottleneck over collision checking.

4 Numerical Studies

In general, what we care about is determined by real life use cases. experimentally confirming the proofs, and characterizing computational complexity. Demonstrating can be used on actual robot.

E.g. num samples vs (failure rate and pathcost

4.1 number samples vs path cost, averaged with error bars, ideally for each of the algs

TODO: discuss expectations about the graphs for each , and the graphs for algs compared to each other

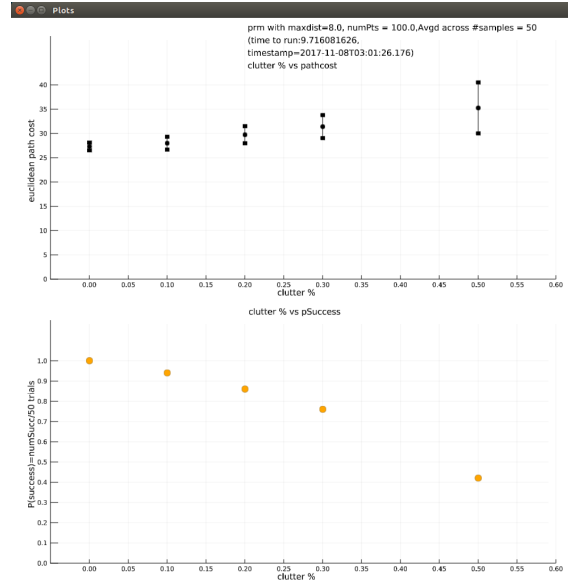


Figure 3: PRM

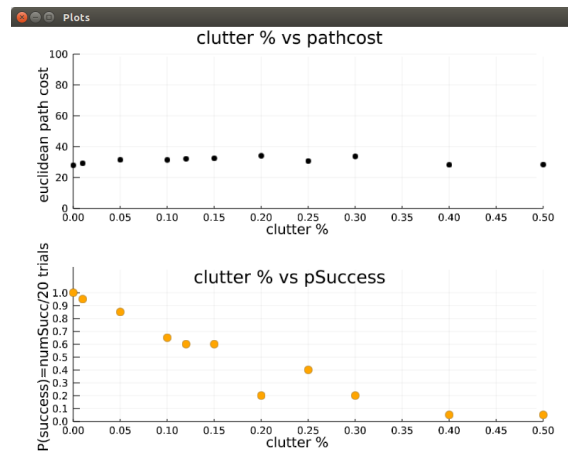


Figure 4: PRM

4.2 clutter vs pathcost

4.3 bugtrap

4.4 Implementation Details

Modularity: Roadmap vs Graph Search vs Plotting individual (and possibly saving plots of individual runs), vs Plotting and collecting data.

Walls: collection of lines

How to deal with graphs so that the information about the actual run (e.g. pathcost, connection radius, etc.) This probably ate 20-30 hours of Julia frustration... I should have used the matplotlib library probably. OR ... just used python. :((e.g. including timestamp very useful).

Node: should be this or that (include cost in definition, but optionally – the optional part ate something like 10 hours in Julia frustration)

Environment: Apparently spiky obstacles are interesting too. As well as higher dimensions.

Maybe document some Julia things... workflow, types system and common errors, where to seek




Figure 5: bugtrap widedoor

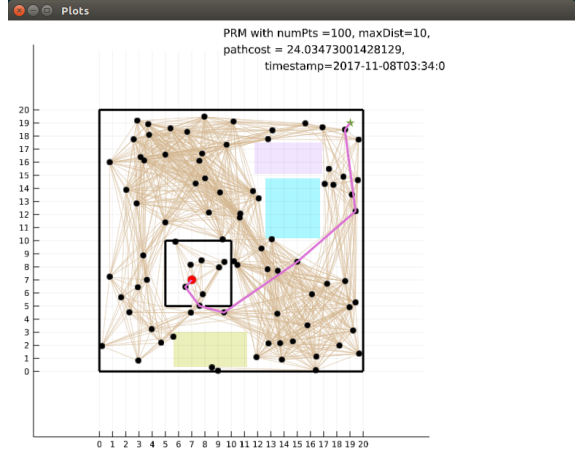


Figure 6: some caption

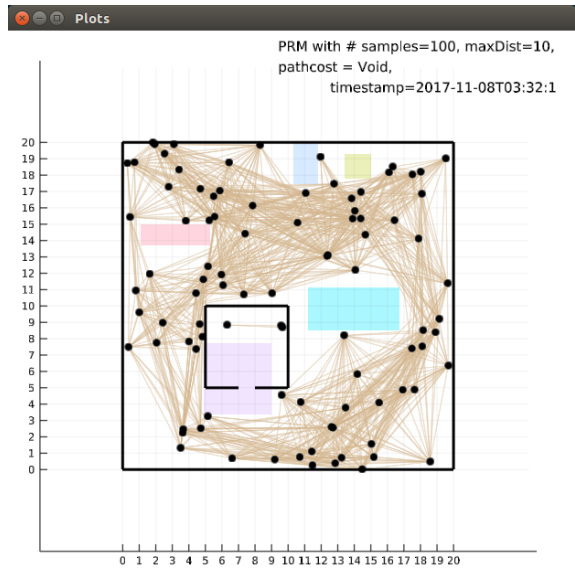


Figure 7: PRM_bugtrap_needObsFeasibilityCheck.png




Figure 8: PRM_bugtrap_doorSize_pSuccess.png

help ⁸ And then discuss pros and cons of Julia and warn people about / away from it.

4.5 Debugging Protips

Trees should not be graphs (no cyclic components)

Lines should not cross obstacles

Plot n = 500 points etc. to see how it covers the room

5 Results

5.1 num samples vs (failure rate and pathcost)

5.1.1 PRM

todo RRT, RRT*, PRM* double check this

5.2 %clutter vs (pathcost and feasibility)

6 Things I learned

I learned about reading math papers. 99% of the symbols are defined in the paper, you just have to be patient about finding the definition either earlier or later in the paper, or perhaps in the appendix somewhere.

⁸(probably also cover how to atone for your sins of trying weird languages not made by large corporations by sacrificing small laptops to the programming gods)

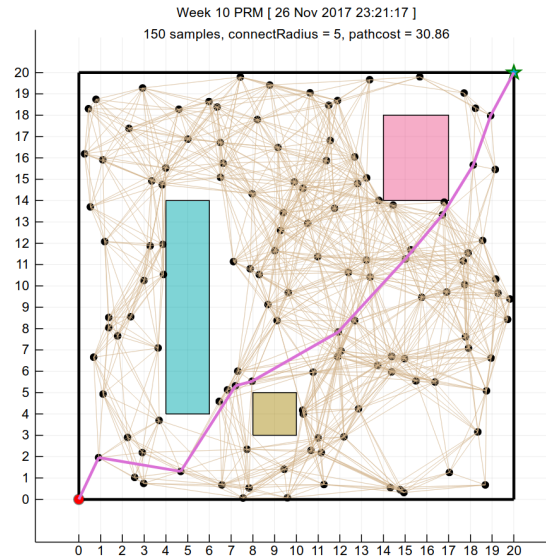


Figure 9: prm many samples

Reading a page of proofs takes >10 hours a page. Should learn to pick battles (understanding some of the stuff may not be critical to understanding the approach of the proof, and also some time should be spent on what the proof tells us / how the proof is useful)

Abstractions – for instance, we don't need to define a robot size or clearance needed, we can simply "increase" the size of the obstacles

<The set minus, the gets, the union and and signs, supremeum and infimum, >

I also learned about the differences between sampling-based motion planning and nonlinear convex optimization approaches, and which one is more suited than the other.

Latex is pretty reasonable, I had some horrific first experience and avoided it since. But just installed all 3 GB of packages or whatever (via `sudo apt install texmaker`) and previewed templates on sharelatex and now all set.

I also learned a bit about keeping in mind the overall goal vs. getting stuck on any particular milestone / to-do item. I don't think I had the chance to do so, but I wish I had been able to make more executive decisions about what to accomplish.

7 Reflection

This rotation was very enjoyable, I've never worked through a proof before. I feel less intimidated by algorithmic notation and proofs now that I have a better idea what they are. In that sense, forcing myself to work through this rotation helped me out a lot. I also hadn't thought math was fun since middle school, so it was really nice to hang out with math majors who were really enthusiastic about math. Hopefully this feeling will keep me going when I take math courses and have to do graded things.

My choice of Julia was a tad unfortunate, although it was kind of neat to see. In retrospect, I really didn't want to deal with MATLAB (which was glitchy on linux, and I didn't have a windows install at the time. I miss public computers), and I thought it might be annoying to understand what was going on with numpy. But in the end I used no matrices... so... python was a much better way to go. I didn't think to ask graph python people (instead of robotics people) to show me a workflow, though a friend helped with the computer vision homework and I got to know the python syntax a lot better.

An example of why Julia documentation and I did not get along: <https://github.com/scheinerman/Permutations.jl>

I think to some extent I expected it to be a little rough to pick to do a PhD in a new-ish field, and then to take both computer vision (linear algebra) and this rotation (statistics) when I hadn't taken a math class in a while. But it was even rougher than I imagined...

Regardless I certainly learned a lot compared to the beginning of the semester! I'm happy, this is part of why I decided to go to grad school. It was neat to see it all tying together, AI and this rotation and even some of the ideas from computer vision.

In the future, I intend to take more math classes and programming classes. I also intend to not use Julia again for a long while. After I switched from MATLAB to python for computer vision final project, it was amazing, I'd forgotten that programming could actually be pleasant and fun, and I imagine something similar will happen with switching from Julia to Python.

8 Conclusion

Math is so cool!

9 Thanks

People who have helped me understand math / Julia:

- Irina Tolikova
- Eric Lu
- Ambarish Chattopadhyay
- Robin Deits
- Misc. patient strangers online who dealt with my deep incoherent frustration... I will look their names up later

Thanks to my advisor Lucas Janson for sparing an hour a week to patiently trying to explain things to me in a way that made sense to me, for writing numerous recommendation letters, dealing with my continual work tardiness, and helping me figure out what I want to do in terms of research.

10 References

References

- [1] LaValle, S. M. (1998). Rapidly-exploring random trees: A new tool for path planning. [Paper link](#)
- [2] Karaman, S., Frazzoli, E. (2011). Sampling-based algorithms for optimal motion planning. The international journal of robotics research, 30(7), 846-894. [PDF link](#)
- [3] Fast Marching Tree *Imperial Japan 1800-1945* 2015. Janson, L., Schmerling, E., Clark, A., Pavone, M. (2015). Fast marching tree: A fast marching sampling-based method for optimal motion planning in many dimensions. The International journal of robotics research, 34(7), 883-921. [Article](#) [DOI link](#) [PDF link](#)